Final Project - Predicting March Madness

Anmol Sapru and Rohit Gunda

Introduction and Data

Motivation

Duke is synonymous with basketball. As Duke students who love Duke Basketball and March Madness, we are interested in performing a statistical analysis on the most thrilling tournament in sports. While watching the 2023 March Madness tournament and the many upsets that came with it, we were motivated to see if we could use statistical methods to predict March Madness winners. Upon scouring sites such as FiveThirtyEight and the KenPom rankings, we were inspired to create models of our own to predict tournament success.

Fundamental Research Question: What variables are important to March Madness success and which outliers over the past 15 years have existed that bring "Madness" to "March?

Packages

```
library(tidyverse)
library(tidymodels)
library(Stat2Data)
library(caret)
library(leaps)
library(MASS)
```

Data

Data was found from the Sports Reference college basketball team stats website and Bart Torvik analytics website, with these general and deeper stats being taken going back to 2008 (excluding the cancelled March Madness of 2020). From here, the data was cleaned in Excel to separate the year and result of each team then joined for all of this data together. With the nature of March Madness being over time, the data most definitely violates expectations

of independence, with similar players, coaches, and more between years, but there would not be a reasonable way to complete such analysis without this independence.

Key Variables

- ADJOE/ADJDE - Points scored per 100 possessions on offense/defense, adjusted for opponent strength and game location
- TOR/TORD - Turnovers committed/forced per game on offense/defense
- ADJT - Estimate possessions per game a team would have against the average tempo
- EFG% - Field goal percentage adjusted for value of baskets scored

Exploratory Data Analysis

EXPLORATORY DATA ANALYSIS-Averages for stats—add by year/by round, Averages for winner stats

Description of Data Cleaning

For further cleaning, the non-March Madness teams added in the join were removed and simple variable names and values were cleaned up to be easier to work with.

ADD CLEARER DESCRIPTION OF DATA CLEANING

Methodology

We determined it would be best to try different regressions, with logistic regressions by round and an ordinal regression—each with their different assumptions to look into. The issue with the first of these is the low levels of data that limit the creation of a model beyond around the Elite Eight. In preparing for the round-by-round logistic regressions we split up the data from the original data sets into multiple data sets for each round, with a TRUE or FALSE value of whether they won their game in that round.

WHY CHOSE EACH VARIABLE

Round-by-Round Logistic Regression

The following is an example of the regression that was ran on all of the Round of 64 teams to create a regression that predicts winners (round_64 = TRUE) against losers (round_64 = FALSE) for the round. Using stepAIC works to attempt to limit the overfitting with the data. This was also done for the Round of 32 and Sweet Sixteen.

WHY LOGISTIC REG?

```
round_64 <- na.omit(round_64)

round_64_max <- glm(round_64 ~ G.x + WINS + LOSSES + ADJOE + ADJDE +
                    `EFG%` + `EFGD%` + TOR + TORD + ORB + DRB + FTR +
                    FTRD + `2P%` + `2P%D` + `3P%` + `3P%D` + `3PR` +
                    `3PRD` + `ADJ T` +`Conf. W-L%` + `Home W-L%` +
                    `Away W-L%` + `AVG PPG` + `AVG DPPG` + `AVG PD` +
                    `AST/TOV` + `PF/G` + WINS*G.x + LOSSES*G.x +
                    ADJOE*ADJDE +`EFG%`*`EFGD%` + TOR*TORD + ORB*DRB +
                    FTR*FTRD + `2P%`*`2P%D` + `3P%`*`3P%D` + `3PR`*`3PRD` +
                    `2P%`*`3P%` + `2P%D`*`3P%D` + `AVG PPG`*`AVG DPPG`,
                    data = round_64,
                    family = "binomial")

round_64_min <- glm(round_64 ~ 1,
                    data = round_64,
                    family = "binomial")

round_64_model <- stepAIC(round_64_max,
        scope = list(lower = round_64_min, upper = round_64_max),
        data = round_64, direction = "both")
```

The following are the regression models that were found from the three rounds of data and regression models run through stepAIC:

Overall Ordinal Regression

```
cbb$march_madness <- fct_relevel(cbb$march_madness,
                          c("R64", "R32", "S16", "E8",
                            "F4", "2ND", "Champions"))

mmb_ord_model <- polr(march_madness ~ ADJOE + ADJDE + TORD +
```

```
                       FTRD + `2P%D` + `3P%D` + `AVG PPG` +
                       `AVG DPPG` + `PF/G`, data = cbb)
```

ADD ASSESSMENT OF CONDITIONS + DIAGNOSTICS

Results


```
`2023sr` <- read_csv("data/2023sportsreference.csv")
```


```
Rows: 363 Columns: 13
-- Column specification --------------------------------------------------------
Delimiter: ","
chr  (1): School
dbl (12): G, Overall W-L%, Overall SRS, Overall SOS, Conf. W-L%, Home W-L%, ...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```


```
`2023analytics` <- read_csv("data/2023torvik.csv")
```


```
Rows: 363 Columns: 22
-- Column specification --------------------------------------------------------
Delimiter: ","
chr  (1): TEAM
dbl (21): RK, G, WINS, LOSSES, ADJOE, ADJDE, EFG%, EFGD%, TOR, TORD, ORB, DR...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```


```
`2023stats` <- left_join(`2023analytics`, `2023sr`, by = c("TEAM" = "School"))

`2023teams` <- read_csv("data/2023teams.csv")
```


```
Rows: 64 Columns: 1
-- Column specification --------------------------------------------------------
Delimiter: ","
```

```
chr (1): TEAM

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

  `2023stats` <- `2023stats` |>
    filter(`2023stats`$TEAM %in% `2023teams`$TEAM)

  tibble(predict(round_64_model, `2023stats`)) |>
    mutate(rank = seq(1:64)) |>
    left_join(mutate(`2023stats`, rank = seq(1:64))) |>
    arrange(desc(predict(round_64_model, `2023stats`)))
```

Joining, by = "rank"

```
# A tibble: 64 x 36
   predi~1  rank    RK TEAM    G.x  WINS LOSSES ADJOE ADJDE `EFG%` `EFGD%`   TOR
     <dbl> <int> <dbl> <chr> <dbl> <dbl>  <dbl> <dbl> <dbl>  <dbl>   <dbl> <dbl>
 1   3.67     1     1 Hous~    34    31      3  117.  88     52.7    42.5  15.3
 2   3.37     3     3 UCLA     34    29      5  113.  87.4   50.9    46.8  15.3
 3   2.54     6     6 Purd~    34    29      5  118.  92.6   52.2    47.2  17
 4   2.48    12    12 San ~    32    26      6  111.  90.1   50.1    47.5  17.6
 5   2.47     2     2 Alab~    34    29      5  115.  88.3   52.7    41.5  19
 6   2.10     9     9 Texas    34    26      8  115.  91.6   52.7    47.8  16.5
 7   2.07    11    11 Marq~    34    28      6  119.  96.1   56      51.1  15.2
 8   2.05     4     4 Tenn~    33    23     10  111.  86.2   50.3    42.4  18.1
 9   1.98    13    13 Kans~    34    27      7  113.  91.5   52.4    47.1  17.5
10   1.85     5     5 Conn~    33    25      8  119.  92.5   53.5    45.5  18.9
# ... with 54 more rows, 24 more variables: TORD <dbl>, ORB <dbl>, DRB <dbl>,
#   FTR <dbl>, FTRD <dbl>, `2P%` <dbl>, `2P%D` <dbl>, `3P%` <dbl>,
#   `3P%D` <dbl>, `3PR` <dbl>, `3PRD` <dbl>, `ADJ T` <dbl>, G.y <dbl>,
#   `Overall W-L%` <dbl>, `Overall SRS` <dbl>, `Overall SOS` <dbl>,
#   `Conf. W-L%` <dbl>, `Home W-L%` <dbl>, `Away W-L%` <dbl>, `AVG PPG` <dbl>,
#   `AVG DPPG` <dbl>, `AVG PD` <dbl>, `AST/TOV` <dbl>, `PF/G` <dbl>, and
#   abbreviated variable name 1: `predict(round_64_model, \`2023stats\`)`
```

```
  tibble(predict(round_32_model, `2023stats`)) |>
    mutate(rank = seq(1:64)) |>
    left_join(mutate(`2023stats`, rank = seq(1:64))) |>
    arrange(desc(predict(round_32_model, `2023stats`)))
```

```
Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
prediction from a rank-deficient fit may be misleading


Joining, by = "rank"


Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
prediction from a rank-deficient fit may be misleading


# A tibble: 64 x 36
   predi~1  rank    RK TEAM     G.x  WINS LOSSES ADJOE ADJDE `EFG%` `EFGD%`   TOR
     <dbl> <int> <dbl> <chr>  <dbl> <dbl>  <dbl> <dbl> <dbl>  <dbl>   <dbl> <dbl>
 1   2.04     1     1 Hous~    34    31      3  117.  88     52.7    42.5  15.3
 2   1.78     3     3 UCLA     34    29      5  113.  87.4   50.9    46.8  15.3
 3   1.67    38    42 Penn~    35    22     13  115.  99.9   55.5    49.2  13.7
 4   1.49     2     2 Alab~    34    29      5  115.  88.3   52.7    41.5  19
 5   1.03    12    12 San ~    32    26      6  111.  90.1   50.1    47.5  17.6
 6   0.996    6     6 Purd~    34    29      5  118.  92.6   52.2    47.2  17
 7   0.750   16    16 Memp~    34    26      8  114.  94.2   53.2    46.8  18.4
 8   0.635   14    14 Crei~    33    21     12  114.  92.8   54.3    47.3  16.6
 9   0.603   10    10 Gonz~    32    27      5  123.  98.6   58.5    51.7  14.6
10   0.574   50    78 Oral~    30    26      4  112. 101.    56.1    48.5  13.2
# ... with 54 more rows, 24 more variables: TORD <dbl>, ORB <dbl>, DRB <dbl>,
#   FTR <dbl>, FTRD <dbl>, `2P%` <dbl>, `2P%D` <dbl>, `3P%` <dbl>,
#   `3P%D` <dbl>, `3PR` <dbl>, `3PRD` <dbl>, `ADJ T` <dbl>, G.y <dbl>,
#   `Overall W-L%` <dbl>, `Overall SRS` <dbl>, `Overall SOS` <dbl>,
#   `Conf. W-L%` <dbl>, `Home W-L%` <dbl>, `Away W-L%` <dbl>, `AVG PPG` <dbl>,
#   `AVG DPPG` <dbl>, `AVG PD` <dbl>, `AST/TOV` <dbl>, `PF/G` <dbl>, and
#   abbreviated variable name 1: `predict(round_32_model, \`2023stats\`)`
```

```r
tibble(predict(sweet_sixteen_model, `2023stats`)) |>
  mutate(rank = seq(1:64)) |>
  left_join(mutate(`2023stats`, rank = seq(1:64))) |>
  arrange(desc(predict(sweet_sixteen_model, `2023stats`)))
```

```
Joining, by = "rank"


# A tibble: 64 x 36
   predi~1  rank    RK TEAM     G.x  WINS LOSSES ADJOE ADJDE `EFG%` `EFGD%`   TOR
     <dbl> <int> <dbl> <chr>  <dbl> <dbl>  <dbl> <dbl> <dbl>  <dbl>   <dbl> <dbl>
 1   2.47     6     6 Purd~    34    29      5  118.  92.6   52.2    47.2  17
```

```
  2    1.84       1       1 Hous~    34    31        3  117.  88     52.7     42.5  15.3
  3    1.30       5       5 Conn~    33    25        8  119.  92.5   53.5     45.5  18.9
  4    1.22      13      13 Kans~    34    27        7  113.  91.5   52.4     47.1  17.5
  5    0.880     22      22 Texa~    34    25        9  113.  94.8   49       47.9  18.3
  6    0.826     24      24 Kans~    32    23        9  110.  93.1   51.5     47.5  20
  7    0.752     21      21 Duke     34    26        8  112.  93.1   51.1     46.3  18.2
  8    0.714     12      12 San ~    32    26        6  111.  90.1   50.1     47.5  17.6
  9    0.657     38      42 Penn~    35    22       13  115.  99.9   55.5     49.2  13.7
 10    0.606     26      26 TCU      33    21       12  110.  93.4   50.2     47.8  17.1
# ... with 54 more rows, 24 more variables: TORD <dbl>, ORB <dbl>, DRB <dbl>,
#   FTR <dbl>, FTRD <dbl>, `2P%` <dbl>, `2P%D` <dbl>, `3P%` <dbl>,
#   `3P%D` <dbl>, `3PR` <dbl>, `3PRD` <dbl>, `ADJ T` <dbl>, G.y <dbl>,
#   `Overall W-L%` <dbl>, `Overall SRS` <dbl>, `Overall SOS` <dbl>,
#   `Conf. W-L%` <dbl>, `Home W-L%` <dbl>, `Away W-L%` <dbl>, `AVG PPG` <dbl>,
#   `AVG DPPG` <dbl>, `AVG PD` <dbl>, `AST/TOV` <dbl>, `PF/G` <dbl>, and
#   abbreviated variable name ...
```

```r
  tibble(predict(mmb_ord_model, `2023stats`)) |>
    mutate(rank = seq(1:64)) |>
    left_join(mutate(`2023stats`, rank = seq(1:64))) |>
    arrange(predict(mmb_ord_model, `2023stats`, type = "probs"))
```

```
Joining, by = "rank"


# A tibble: 64 x 36
   predi~1  rank     RK TEAM    G.x  WINS LOSSES ADJOE ADJDE `EFG%` `EFGD%`   TOR
   <fct>   <int>  <dbl> <chr> <dbl> <dbl>  <dbl> <dbl> <dbl>  <dbl>   <dbl> <dbl>
 1 E8          1      1 Hous~    34    31      3  117.  88     52.7    42.5  15.3
 2 S16         3      3 UCLA     34    29      5  113.  87.4   50.9    46.8  15.3
 3 S16         5      5 Conn~    33    25      8  119.  92.5   53.5    45.5  18.9
 4 S16         4      4 Tenn~    33    23     10  111.  86.2   50.3    42.4  18.1
 5 R32         6      6 Purd~    34    29      5  118.  92.6   52.2    47.2  17
 6 R32         2      2 Alab~    34    29      5  115.  88.3   52.7    41.5  19
 7 R32         9      9 Texas    34    26      8  115.  91.6   52.7    47.8  16.5
 8 R32        11     11 Marq~    34    28      6  119.  96.1   56      51.1  15.2
 9 R32        10     10 Gonz~    32    27      5  123.  98.6   58.5    51.7  14.6
10 R32         8      8 Sain~    32    25      7  112.  89.1   52.5    46.7  16.4
# ... with 54 more rows, 24 more variables: TORD <dbl>, ORB <dbl>, DRB <dbl>,
#   FTR <dbl>, FTRD <dbl>, `2P%` <dbl>, `2P%D` <dbl>, `3P%` <dbl>,
#   `3P%D` <dbl>, `3PR` <dbl>, `3PRD` <dbl>, `ADJ T` <dbl>, G.y <dbl>,
#   `Overall W-L%` <dbl>, `Overall SRS` <dbl>, `Overall SOS` <dbl>,
```

```
#   `Conf. W-L%` <dbl>, `Home W-L%` <dbl>, `Away W-L%` <dbl>, `AVG PPG` <dbl>,
#   `AVG DPPG` <dbl>, `AVG PD` <dbl>, `AST/TOV` <dbl>, `PF/G` <dbl>, and
#   abbreviated variable name 1: `predict(mmb_ord_model, \`2023stats\`)`
```

INTERPRET SOMETHING - MOST IMPORTANT VARIABLES

ADD ASSUMPTIONS

MODEL'S PREDICTIVE POWER

Outliers and Randomness of March Madness

Discussion

MISSING ASSUMPTIONS

HYPOTHESIS TESTS

COOK'S DISTANCE

IDEA'S FOR FUTURE WORK